

Chain of Responsibility Pattern

Introduction

The Chain of Responsibility Design Pattern is a behavioral design pattern that allows you to pass requests along a chain of handlers until one of them handles the request. Imagine you have a customer support system where requests come in and need to be handled at different levels i.e. first by a support agent, then by a supervisor if the agent cannot handle it, then by a manager if the supervisor cannot handle it, and finally by a director if the manager cannot handle it. Without the Chain of Responsibility pattern, you would need to create complex conditional logic with nested if-else statements or switch cases to determine which handler should process each request. This means modifying the request handling logic every time you need to add a new handler level or change the handling rules, which violates the Open/Closed Principle. The Chain of Responsibility pattern solves this by creating a chain of handler objects where each handler has a reference to the next handler in the chain. When a request comes in, it is passed along the chain until a handler can process it. Just like how a customer service request escalates through different levels of support staff until someone can resolve it, a request in the Chain of Responsibility pattern moves through different handler objects until one can handle it. The pattern is called Chain of Responsibility because it creates a chain where each link has the responsibility to either handle the request or pass it to the next link in the chain.

Problem Statement

Consider a scenario where you have a customer support system with different handler levels such as SupportAgent, Supervisor, Manager, and Director. Your application needs to process various types of requests like technical issues, billing inquiries, refund requests, complaints, and escalation requests. Without the Chain of Responsibility pattern, you would need to create complex conditional logic with nested if-else statements or switch cases to determine which handler should process each request based on its type, priority, or complexity. This means if you have 4 handler levels and 5 request types, you must implement complex decision logic with multiple conditions checking request type, priority thresholds, and handler capabilities. With 6 handler levels and 10 request types you need even more complex conditional logic with dozens of nested conditions. Without Chain of Responsibility Pattern you need to create separate methods like handleTechnicalRequest(), handleBillingRequest(), handleRefundRequest(), or use massive switch statements with nested if-else blocks checking request priority and handler availability. This approach leads to massive code

duplication because similar decision logic is repeated across multiple request handling methods. For example all request handling methods will have similar structure for checking handler capabilities, priority thresholds, and escalation rules, even though the actual decision criteria differs for each request type. When you want to add a new handler level like TeamLead, you must modify all request handling methods by adding new conditions and priority checks. Similarly adding a new request type requires implementing decision logic in all existing handler methods.

This tight coupling between request types and handler levels makes the system very difficult to extend and maintain. Every time you add a new handler level, you must touch and modify all request handling logic in your system. If you need to change how a request is routed or which handler can process it, you must update that logic in multiple places. This violates the Open/Closed Principle because classes should be open for extension but closed for modification. The request handling classes become bloated with decision logic that mixes concerns - the handler class should represent the handler's capabilities and responsibilities, not all possible routing and decision logic for every request type. If routing logic needs to access private members of handler classes, you must break encapsulation by exposing internal details. The classes become overly complex with too many responsibilities, violating the Single Responsibility Principle. The conditional logic becomes hard to read, test, and maintain as it grows with more handlers and request types, making the system rigid and fragile.

Solution

The Chain of Responsibility Design Pattern solves this problem by separating the request handling logic from the routing decision logic. Instead of using complex conditional statements to determine which handler should process each request, it creates a chain of handler objects where each handler has a reference to the next handler in the chain. One hierarchy is created for handlers i.e. SupportAgent, Supervisor, Manager and Director, all implementing a common Handler interface that defines a handleRequest(Request) method and a setNext(Handler) method. Another separate abstraction is created for requests i.e. TechnicalRequest, BillingRequest, RefundRequest, each containing the request data and type information.

The handler classes don't directly implement complex routing logic. Instead each handler class only implements the handleRequest(Request) method that checks if it can process the request. If the handler can process the request, it does so and stops. If it cannot handle the request, it passes the request to the next handler in the chain by calling nextHandler.handleRequest(request). The chain is constructed by linking handlers together using the setNext() method. For example,

`SupportAgent.setNext(Supervisor)`, `Supervisor.setNext(Manager)`, `Manager.setNext(Director)`. When a request comes in, it is passed to the first handler in the chain. The request then moves through the chain until a handler can process it or the chain ends.

This separation allows adding new handlers by creating new handler classes and inserting them into the chain without modifying existing handlers, and adding new request types by creating new request classes without touching handler classes. The pattern uses delegation - each handler delegates to the next handler if it cannot process the request. This loose coupling makes the system easy to extend and maintain. Changes to handler behavior only require updating the specific handler class, and changes to request structure only affect the request classes. The chain can be dynamically reconfigured at runtime, allowing handlers to be added, removed, or reordered without modifying existing code. Related handlers are grouped together in a chain making the code more organized and maintainable.

Code

Code implementation of Chain of Responsibility pattern:

<https://github.com/prateek27/design-patterns-java/tree/main/design-patterns/src/main/java/org/prateek/BehaviouralPatterns/ChainOfResponsibilityPattern>

Use Cases

The Chain of Responsibility Design Pattern is most useful when you have a stable set of handlers but frequently need to add new handler levels or change the order of handlers. If the handler chain structure changes rarely but handlers are added or reordered frequently, Chain of Responsibility pattern provides the perfect solution.

Chain of Responsibility pattern should be used when there is a need to route requests to different handlers without using complex conditional logic. This is particularly important when request routing needs to be flexible and handlers need to be added without changing existing code, following the Open/Closed Principle. Consider a customer support system where handler levels remain stable but new handlers like TeamLead or SeniorManager are added frequently, or the chain order needs to be changed based on business requirements.

The pattern is highly beneficial when you want to separate routing logic from handler implementation. If multiple handlers need to work together in a sequence, Chain of

Responsibility pattern allows grouping all handlers in a chain, making the code more organized and maintainable. This is useful when handlers need to share state or work together in a specific order.

Chain of Responsibility pattern works well when you need dynamic chain configuration at runtime. Instead of hardcoding routing logic with if-else statements, the pattern allows handlers to be added, removed, or reordered dynamically without modifying existing code. The pattern is also suitable when you want to avoid complex conditional logic when processing requests that need to be handled by different handlers based on their type, priority, or complexity.

However, the Chain of Responsibility pattern should not be used when the handler chain structure changes frequently. If handlers are constantly being added, removed, or reordered, the overhead of maintaining the chain structure may not be justified. The pattern is not suitable for systems where the handler chain is unstable or changes frequently.

The pattern is also not appropriate when there are only a few handlers with simple routing logic. If the routing logic is straightforward and unlikely to change, the overhead of maintaining handler classes and the chain mechanism is not justified. Simple conditional statements would be more straightforward. The pattern should only be used when there is genuine need for flexible request routing with multiple handlers that may change over time.

Applications

Customer Support Systems: Customer support platforms use Chain of Responsibility pattern extensively for request escalation. The handler chain with levels like support agent, supervisor, manager, and director remains stable, while new handler levels like team lead or senior manager are added frequently. Handlers like `SupportAgentHandler`, `SupervisorHandler`, `ManagerHandler`, and `DirectorHandler` process requests and pass them to the next handler if they cannot handle them. Applications like help desk systems, ticketing systems, and customer service platforms use this pattern.

Web Request Processing: Web frameworks use Chain of Responsibility pattern for middleware chains and request processing pipelines. The request processing chain with handlers like authentication, authorization, validation, logging, and error handling remains stable, while new middleware handlers are added frequently. Handlers like `AuthenticationHandler`, `AuthorizationHandler`, `ValidationHandler`, and `LoggingHandler`

process HTTP requests and pass them to the next handler in the chain. Frameworks like Express.js, Spring MVC, ASP.NET Core, and Django use this pattern.

Exception Handling: Programming languages and frameworks use Chain of Responsibility pattern for exception handling. The exception handling chain with handlers for different exception types remains stable, while new exception handlers are added frequently. Handlers like NullPointerExceptionHandler, IllegalArgumentExceptionHandler, and IOExceptionHandler catch exceptions and pass them to the next handler if they cannot handle them. Languages like Java, C#, Python, and their exception handling mechanisms use this pattern.

Logging Systems: Logging frameworks use Chain of Responsibility pattern for log level filtering and processing. The logging chain with handlers for different log levels like debug, info, warning, and error remains stable, while new log handlers or filters are added frequently. Handlers like DebugHandler, InfoHandler, WarningHandler, and ErrorHandler process log messages and pass them to the next handler based on log level. Frameworks like Log4j, SLF4J, Winston, and Python logging use this pattern.

Advantages and Disadvantages

Advantages

Chain of Responsibility pattern provides strong separation between request handling logic and routing decision logic. This separation allows both to be developed, tested and maintained independently. Changes in handler behavior do not affect the routing logic and vice versa, making the codebase more stable and easier to manage.

The pattern enables easy addition of new handlers without modifying existing handlers. New handlers can be added by creating new handler classes and inserting them into the chain without touching existing handler code, following the Open/Closed Principle where classes are open for extension but closed for modification. This allows the system to evolve by adding new functionality without breaking existing code.

Chain of Responsibility pattern groups related handlers together in a chain, making the code more organized and maintainable. All handlers that work together in a sequence are contained in one chain, making it easier to understand, test and modify handler behavior. Handlers can also share state and work together within a chain.

The pattern allows dynamic chain configuration at runtime. Instead of hardcoding routing logic with if-else statements, handlers can be added, removed, or reordered

dynamically without modifying existing code. This provides flexibility that static conditional logic cannot provide.

Chain of Responsibility pattern eliminates the need for complex conditional logic when routing requests to different handlers. The chain mechanism automatically routes requests through handlers until one can process them, making code cleaner and less error-prone. This improves code readability and maintainability.

The pattern improves code reusability as handlers can be shared across multiple chains. It also makes handlers testable independently since each handler can be tested in isolation without requiring the full chain structure. The chain can be easily reconfigured for different scenarios or requirements.

Disadvantages

Chain of Responsibility pattern increases overall complexity of the codebase. It introduces additional classes, interfaces and the chain mechanism which can make the system harder to understand initially. For simple scenarios, this added complexity may not be justified.

The pattern adds a level of indirection through the chain mechanism. This can make the code slightly harder to follow and debug as method calls pass through multiple handlers in the chain. There is also a minor performance overhead due to this indirection, especially if requests pass through many handlers before being processed.

Chain of Responsibility pattern requires careful upfront design to identify the right handler chain structure and order. If the chain structure is not stable or if handlers are unlikely to be added, the pattern may not provide the intended benefits and could make the system more complicated than necessary.

When new handlers are added, you must carefully configure the chain to insert them in the correct position. If the chain order is critical, this can be error-prone and requires careful management. If handlers are added or removed frequently, this overhead becomes significant.

For systems with only a few handlers or simple routing logic, Chain of Responsibility pattern is overkill. The overhead of maintaining handler classes, the Handler interface, and the chain mechanism is not worth it when simpler approaches like conditional statements would suffice. The pattern should only be used when there is genuine need for flexible request routing with multiple handlers that may change over time.